



Uni.lu HPC School 2021

PS10b: Introduction to OpenCL Programming

Uni.lu High Performance Computing (HPC) Team
T. Carneiro L. Koutsantonis

University of Luxembourg (UL), Luxembourg

High Performance
Computing &
Big Data Services

 hpc.uni.lu

 hpc@uni.lu

 @ULHPC

 **EMBOURG**
LET'S MAKE IT HAPPEN

<https://hpc.uni.lu/>

Introduction: what's OpenCL?

- In **OpenACC** the compiler is responsible for the parallelization...
 - which might not work in all scenarios
 - may not result in the best performance
 - It is not portable

In turn...

- **OpenCL** is a standard for heterogeneous programming
 - multicore CPUs, GPUs (AMD, Intel, ARM),
 - FPGAs, Apple M1, tensor cores, and ARM
- **with minor or no modifications.**

Introduction: portability

- OpenCL vs. CUDA - Portability (**)



VS.

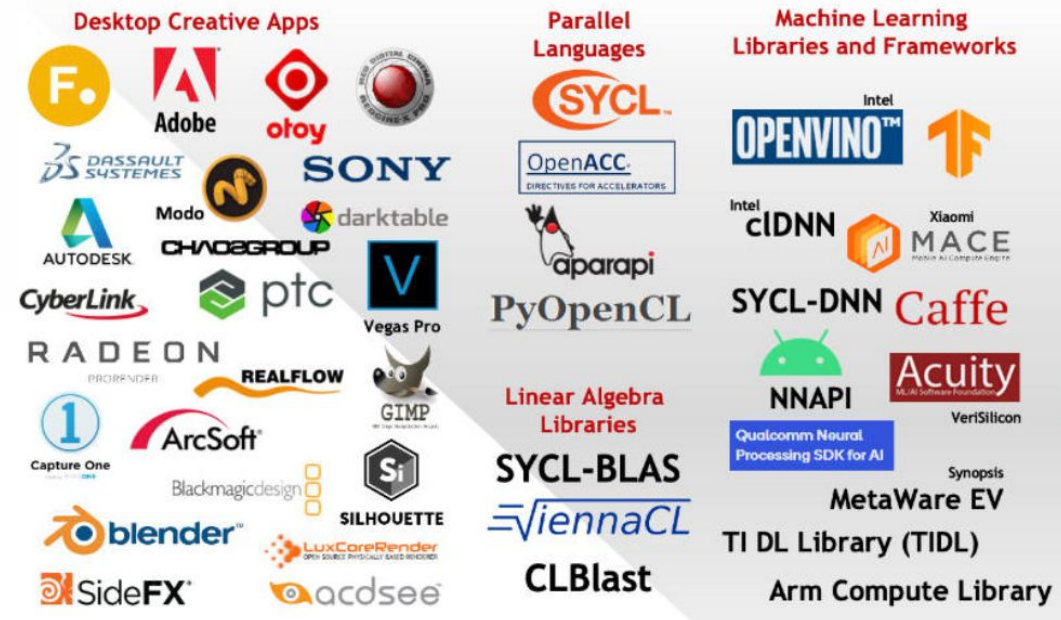


Images from:

<https://www.khronos.org/opencl/>

Introduction: target audience

- **Target audience of OpenCL:** programmers that aim at programming portable heterogeneous code and that want full control of the parallelization process.
 - **Portability has a price:**
OpenCL is much lower level than OpenACC and even then CUDA.



Images from:

<https://www.khronos.org/opencl/>

Have components/ are programmed in OpenCL



Introduction: objectives

- In this tutorial, we teach how to perform the sum of two vectors $C=A+B$ in OpenCL
- **Objectives of the tutorial:**
 - The **main objective** of this tutorial is to introduce for students of the HPC school the heterogeneous programming standard - OpenCL
- **This tutorial covers:**
 - Check for OpenCL-capable device(s);
 - Memory allocation on the device;
 - Data transfer to the device;
 - Retrieve data from the device;
 - Compile C/C++ programs that launch OpenCL kernels.

The OpenCL Platform Model

- **Similar to the CUDA Programming model:**
 - **The code has two parts:** the host and the device parts.
 - **According to the OpenCL specification:** "The model consists of a host (usually the CPU) connected to one or more OpenCL devices (e.g., GPUs, FPGAs)."





Image from the OpenCL specification
3.0.9

The OpenCL Platform Model

- **Similar to the CUDA Programming model:**
 - **The code has two parts:** the host and the device parts.
 - **According to the OpenCL specification:** "The model consists of a host (usually the CPU) connected to one or more OpenCL devices (e.g., GPUs, FPGAs)."



Host: launches the code that is executed on the device(s).

Image from the OpenCL specification
3.0.9

The OpenCL Platform Memory Model

- **Similar to the CUDA programming model:**
 - the host communicates with the *device(s)* through the global memory of the device(s).
 - As in the CUDA programming model, there is a **memory hierarchy** on the device.




Image from Mattson, T.,
McIntosh-Smith, S., Koniges, A.
OpenCL: a Hands-on Introduction

Starting an OpenCL Context

- Initially, OpenCL needs to initialize a context
 - Contexts are used by the OpenCL runtime for *managing* command queues, memory, and for executing kernels on one or more devices connected to the context (*The OpenCL Cookbook*).

```
cl::Context context(CL_DEVICE_TYPE_DEFAULT);
```




Image from Mattson, T., McIntosh-Smith, S., Koniges, A. *OpenCL: a Hands-on Introduction*

Starting an OpenCL Context

- Initially, OpenCL needs to initialize a context
 - Contexts are used by the OpenCL runtime for *managing* command queues, memory, and for executing kernels on one or more devices connected to the context (*The OpenCL Cookbook*).

```
cl::Context context(CL_DEVICE_TYPE_DEFAULT);
```




Image from Mattson, T., McIntosh-Smith, S., Koniges, A. *OpenCL: a Hands-on Introduction*

Starting an OpenCL Context

- **Initially, OpenCL needs to initialize a context**
 - Contexts are used by the OpenCL runtime for *managing* command queues, memory, and for executing kernels on one or more devices connected to the the context [0].

```
cl::Context context(CL_DEVICE_TYPE_DEFAULT);
```

- **Then, it is required to initialize a queue:**
 - Push commands to the device.
 - For those who program in CUDA, OpenCL queues similar to CUDA streams.

```
cl::CommandQueue queue(context);
```




Image from Mattson, T., McIntosh-Smith, S., Koniges, A. *OpenCL: a Hands-on Introduction*

Allocating Memory on the Device: cl::Buffer

- As the device works with vectors A,B,C, we create three buffers:

```
cl::Buffer A_d(context, CL_MEM_READ_ONLY, sizeof(int) * SIZE);  
cl::Buffer B_d(context, CL_MEM_READ_ONLY, sizeof(int) * SIZE);  
cl::Buffer C_d(context, CL_MEM_WRITE_ONLY, sizeof(int) * SIZE);
```

- Using the queue, we write to the buffers of A and B to send their values of A and B from the host to the device:

```
queue.enqueueWriteBuffer(A_d, CL_TRUE, 0, sizeof(int) * SIZE, A_h);  
queue.enqueueWriteBuffer(B_d, CL_TRUE, 0, sizeof(int) * SIZE, B_h);
```

Online Compilation of the Kernel

- For the sake of greater portability, the kernel is provided as a string:

```
std::string kernel_code =  
    "    void kernel_simple_add(global const int* A, global const int* B, global int* C){ "  
    "        C[get_global_id(0)]=A[get_global_id(0)]+B[get_global_id(0)];        "  
    "    }"
```

- Kernels need to return the *void* type.
- The *global* keyword means that the buffer lies on the *global memory*

Building the Kernel

- To build the kernel, we append the kernel string to the program source:

```
cl::Program::Sources sources;  
sources.push_back({ kernel_code.c_str(), kernel_code.length() });
```

- Then, we create a program object using the program source as argument:

```
cl::Program program(context, sources);
```

- **Important:** in this tutorial we use the online compilation of the kernel. Thus, compilation errors of the *device* code are known in **execution time**.

```
if (program.build({ default_device }) != CL_SUCCESS)
```

Building the Kernel

- From the *program*, which contains the *simple_add* kernel, we instantiate a kernel object for execution with three *cl:buffers* as arguments:

```
cl::make_kernel<cl::Buffer, cl::Buffer, cl::Buffer> simple_add(cl::Kernel(program, "simple_add"));
```

- Now, we are able to invoke the kernel *simple_add*:

```
cl::NDRange global(SIZE);  
simple_add(cl::EnqueueArgs(queue, global), A_d, B_d, C_d).wait();
```

- And retrieve the result by using the *queue*:

```
queue.enqueueReadBuffer(buffer_C, CL_TRUE, 0, sizeof(int) * SIZE, C);
```

Compiling the Code

- It is simple to compile an OpenCL code with *gcc* or *g++*. In short, it is only required to add `-lOpenCL` flag to the compilation command.

```
$ g++ -lOpenCL exercise1.cpp
```

- **Exercise 1:** what's the output of the first exercise? Is the default OpenCL platform of Iris the NVIDIA platform?
- **Exercise 2:** following the online tutorial, compile *exercise1.cpp* using the NVIDIA implementation of OpenCL (`-I` and `-L`).
- What's the OpenCL platform after compiling *exercise1.cpp* using the NVIDIA implementation of OpenCL?
- **Obs:** *solutions are in the online tutorial.*

Conclusions

- OpenCL is an **open standard** for heterogeneous programming.
 - However, it exposes the programmer to **much lower-level** concepts than OpenACC and CUDA.
- Depending on the application, one might get a similar performance compared to OpenCL or CUDA using OpenACC.
- **Thus, the programmer needs to decide whether to follow a lower level approach:**
 - Need for portability - OpenCL;
 - The OpenACC compiler can't accelerate the code properly;
 - The OpenACC code became too complex and it is worth to move to OpenCL or CUDA.



Thank you for your attention!



University of Luxembourg, Belval Campus
Maison du Nombre, 4th floor
2, avenue de l'Université
L-4365 Esch-sur-Alzette
mail: hpc@uni.lu



<https://hpc.uni.lu/>



References

- [The OpenCL Cookbook](#)
- [Tutorial: Simple start with OpenCL and C++](#)
- [Khronos OpenCL Working Group. The OpenCL Specification \(Oct. 2021\)](#)
- [Smistad, E. Getting started with OpenCL and GPU Computing, Feb. 22, 2018 \(Access on Oct. 28, 2021\).](#)
- [Mattson, T., McIntosh-Smith, S., Koniges, A. OpenCL: a Hands-on Introduction](#)