



# Uni.lu HPC School 2021

## PS4: HPC Management of Sequential and Embarrassingly Parallel Jobs

High Performance  
Computing &  
Big Data Services



Uni.lu High Performance Computing (HPC) Team  
Dr. S. Varrette

University of Luxembourg (UL), Luxembourg  
<http://hpc.uni.lu>



## Latest versions available on Github:



UL HPC tutorials:

<https://github.com/ULHPC/tutorials>

UL HPC School:

[hpc.uni.lu/education/hpcschool](https://hpc.uni.lu/education/hpcschool)

PS4 tutorial sources:

[ulhpc-tutorials.rtf.d.io/en/latest/sequential/basics](https://ulhpc-tutorials.rtf.d.io/en/latest/sequential/basics)





# Summary

- 1 Introduction**  
[Serial] Task, Job and scheduling considerations  
Slurm Launchers Templates for Sequential Applications
- 2 A First launcher (1 job, 1 task on one core)
- 3 Embarrassingly parallel jobs with GNU Parallel launcher (1 job, 1 node, N tasks)

## Main Objectives of this Session

- Understanding the **jobs and tasks concepts**
- Recognizing **embarrassingly parallel tasks**
  - ↳ Understanding scheduling challenges induced by their execution even on an HPC facility
- Guidelines to the **optimized** design of HPC job campaign in such case
  - ↳ prepare `run_*` task script
  - ↳ **launcher.\* design** for tasklist management, exploiting Slurm & HW characteristics
  - ↳ **interactive** tests and performance analysis
  - ↳ **passive runs** and speedup analysis compared to seq. executions **up to 92% improvement**
- Discovering GNU Parallel
  - ↳ usage and exploitation of ULHPC Generic launcher for GNU parallel

### Test Case for the session

- Sample “Stress Me!” parameter exploration job campaign
  - ↳ `run_stressme <N>` imposes configurable amount of *stress* on 1 core for `<N>` seconds

# Jobs, Tasks & Local Execution



```
$> ./myprog
```



# Jobs, Tasks & Local Execution



```
$> ./myprog
```



## Jobs, Tasks & Local Execution



```
$> ./myprog  
$> ./myprog -n 10
```



## Jobs, Tasks & Local Execution



```
$> ./myprog  
$> ./myprog -n 10
```





## Jobs, Tasks & Local Execution



```
$> ./myprog  
$> ./myprog -n 10  
$> ./myprog -n 100
```



## Jobs, Tasks & Local Execution



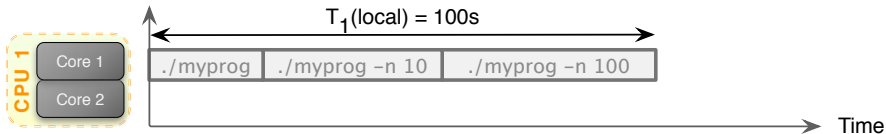
```
$> ./myprog  
$> ./myprog -n 10  
$> ./myprog -n 100
```



## Jobs, Tasks & Local Execution



```
$> ./myprog  
$> ./myprog -n 10  
$> ./myprog -n 100
```



# Jobs, Tasks & Local Execution




**Job(s)**  
 \$> ./myprog  
 \$> ./myprog -n 10  
 \$> ./myprog -n 100

3

**Task(s)**

3



# Jobs, Tasks & Local Execution



```
# launcher  
  
./myprog  
./myprog -n 10  
./myprog -n 100
```



## Jobs, Tasks & Local Execution



```
# launcher  
  
./myprog  
./myprog -n 10  
./myprog -n 100
```



# Jobs, Tasks & Local Execution



```
# launcher  
  
./myprog  
./myprog -n 10  
./myprog -n 100
```



## Jobs, Tasks & Local Execution



```
# launcher  
  
./myprog  
./myprog -n 10  
./myprog -n 100
```






## Jobs, Tasks & Local Execution



```
# launcher  
  
./myprog  
./myprog -n 10  
./myprog -n 100
```




# Jobs, Tasks & Local Execution



```
# launcher
./myprog
./myprog -n 10
./myprog -n 100
```

(s)Job 1

Task(s) 3



# Jobs, Tasks & Local Execution



```
# launcher  
  
./myprog  
./myprog -n 10  
./myprog -n 100
```



# Jobs, Tasks & Local Execution



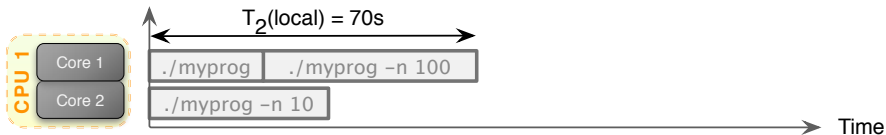
```
# launcher2
"Run in //:"
./myprog
./myprog -n 10
./myprog -n 100
```



# Jobs, Tasks & Local Execution



```
# launcher2
"Run in //:"
./myprog
./myprog -n 10
./myprog -n 100
```



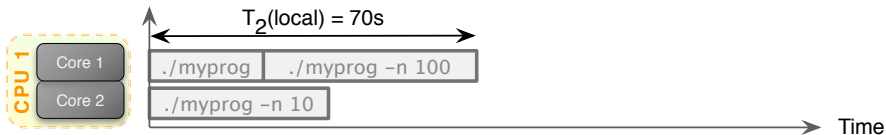
# Jobs, Tasks & Local Execution



```
# launcher2
"Run in //:"
./myprog
./myprog -n 10
./myprog -n 100
```

(s)job 1

Task(s) 3




# Jobs, Tasks & HPC Execution



```
# launcher
./myprog
./myprog -n 10
./myprog -n 100
```

- Note: Simplified example! in practice:

- iris nodes (28c): 2 CPUs of 14 cores each
- aion nodes (128c): (virtual) 8 CPUs of 16 cores each

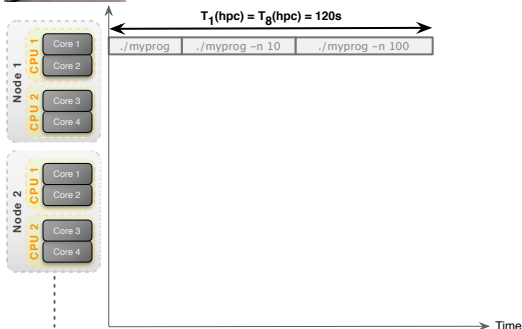


## Jobs, Tasks & HPC Execution



```
# launcher
./myprog
./myprog -n 10
./myprog -n 100
```

- Note: Simplified example! in practice:
  - iris nodes (28c): 2 CPUs of 14 cores each
  - aion nodes (128c): (virtual) 8 CPUs of 16 cores each





## Jobs, Tasks & HPC Execution




```
# launcher
./myprog
./myprog -n 10
./myprog -n 100
```

Job(s) 1

Task(s) 3

- Note: Simplified example! in practice:

- iris nodes (28c): 2 CPUs of 14 cores each
- aion nodes (128c): (virtual) 8 CPUs of 16 cores each




# Jobs, Tasks & HPC Execution



```
# launcher2
"Run in //:"
./myprog
./myprog -n 10
./myprog -n 100
```

- Note: Simplified example! in practice:

- iris nodes (28c): 2 CPUs of 14 cores each
- aion nodes (128c): (virtual) 8 CPUs of 16 cores each



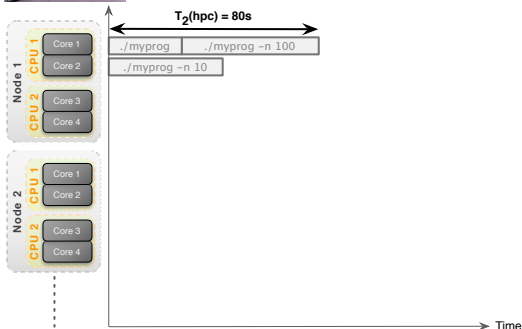
## Jobs, Tasks & HPC Execution



```
# launcher2
"Run in //:"
./myprog
./myprog -n 10
./myprog -n 100
```

- Note: Simplified example! in practice:

- iris nodes (28c): 2 CPUs of 14 cores each
- aion nodes (128c): (virtual) 8 CPUs of 16 cores each



# Jobs, Tasks & HPC Execution



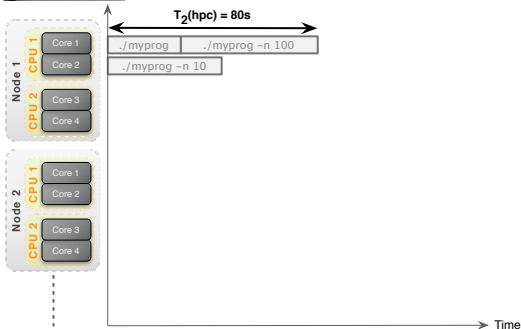
```
# launcher2
"Run in //:"
./myprog
./myprog -n 10
./myprog -n 100
```

Job(s) 1

Task(s) 3

- Note: Simplified example! in practice:

- iris nodes (28c): 2 CPUs of 14 cores each
- aion nodes (128c): (virtual) 8 CPUs of 16 cores each



## Jobs, Tasks & HPC Execution



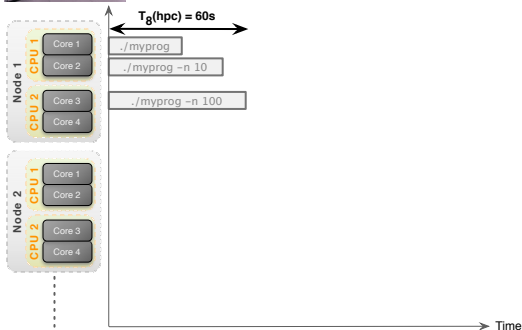
```
# launcher2
"Run in //:"
./myprog
./myprog -n 10
./myprog -n 100
```

Job(s) 1

Task(s) 3

- Note: Simplified example! in practice:

- iris nodes (28c): 2 CPUs of 14 cores each
- aion nodes (128c): (virtual) 8 CPUs of 16 cores each



## Local vs. HPC Executions

Context	Local PC	HPC
Sequential	$T_1(\text{local}) = 100$	$T_1(\text{hpc}) = 120\text{s}$
Parallel/Distributed	$T_2(\text{local}) = 70\text{s}$	$T_2(\text{hpc}) = 80\text{s}$ $T_8(\text{hpc}) = 60\text{s}$

## Local vs. HPC Executions

Context	Local PC	HPC
Sequential	$T_1(\text{local}) = 100$	$T_1(\text{hpc}) = 120\text{s}$
Parallel/Distributed	$T_2(\text{local}) = 70\text{s}$	$T_2(\text{hpc}) = 80\text{s}$ $T_8(\text{hpc}) = 60\text{s}$

- Sequential runs **WON'T BE FASTER** on HPC  
 ↪ Reason: Processor Frequency (typically  $\geq 3\text{GHz}$  vs  $\geq 2\text{GHz}$ )

## Local vs. HPC Executions

Context	Local PC	HPC
Sequential	$T_1(\text{local}) = 100$	$T_1(\text{hpc}) = 120\text{s}$
Parallel/Distributed	$T_2(\text{local}) = 70\text{s}$	$T_2(\text{hpc}) = 80\text{s}$ $T_8(\text{hpc}) = 60\text{s}$

- Sequential runs **WON'T BE FASTER** on HPC
  - Reason: Processor Frequency (typically  $\geq 3\text{GHz}$  vs  $\geq 2\text{GHz}$ )
- Parallel/Distributed runs **DO NOT COME FOR FREE**
  - runs **will be sequential** even if you reserve  $\geq 2$  cores/nodes
  - you have to **explicitly** adapt your jobs to benefit from the multi-cores/nodes



## Serial Task script Launcher

```
#!/bin/bash -l      # <--- DO NOT FORGET '-l'
#SBATCH -N 1
#SBATCH --ntasks-per-node=1
#SBATCH -c 1
#SBATCH --time=0-01:00:00
#SBATCH -p batch
print_error_and_exit() { echo "***ERROR*** $*"; exit 1; }
module purge || print_error_and_exit "No 'module' command"
# C/C++: module load toolchain/intel # OR: module load toolchain/foss
# Java: module load lang/Java/1.8
# Ruby/Perl/Rust...: module load lang/{Ruby,Perl,Rust...}
# /\ ADAPT TASK variable accordingly - absolute path to the (serial) task to be executed
TASK=${TASK:=${HOME}/bin/app.exe}
OPTS=$*

${TASK} ${OPTS}
```

# Serial Python Slurm Launcher

```
#!/bin/bash -l
#SBATCH -N 1
#SBATCH --ntasks-per-node=1
#SBATCH -c 1
#SBATCH --time=0-01:00:00
#SBATCH -p batch
print_error_and_exit() { echo "***ERROR*** $*"; exit 1; }
module purge || print_error_and_exit "No 'module' command"
# Python 3.X by default (also on system)
module load lang/Python
# module load lang/SciPy-bundle
# and/or: activate the virtualenv <name> you previously generated with
#     python -m venv <name>
source ./<name>/bin/activate

python [...]
```

## R Slurm Launcher

```
#!/bin/bash -l
#SBATCH -N 1
#SBATCH --ntasks-per-node=1
#SBATCH -c 28
#SBATCH --time=0-01:00:00
#SBATCH -p batch

print_error_and_exit() { echo "***ERROR*** $*"; exit 1; }
module purge || print_error_and_exit "No 'module' command"
module load lang/R
export OMP_NUM_THREADS=${SLURM_CPUS_PER_TASK:-1}

Rscript <script>.R |& tee job_${SLURM_JOB_NAME}.out
```



# MATLAB Slurm Launcher

```
#!/bin/bash -l
#SBATCH -N 1
#SBATCH --ntasks-per-node=1
#SBATCH -c 28
#SBATCH --time=0-01:00:00
#SBATCH -p batch

print_error_and_exit() { echo "***ERROR*** $*"; exit 1; }
module purge || print_error_and_exit "No 'module' command"
module load base/MATLAB

matlab -nodisplay -nosplash < INPUTFILE.m > OUTPUTFILE.out
```



A First launcher (1 job, 1 task on one core)

---

## Summary

- 1 Introduction  
[Serial] Task, Job and scheduling considerations  
Slurm Launchers Templates for Sequential Applications
- 2 A First launcher (1 job, 1 task on one core)**
- 3 Embarrassingly parallel jobs with GNU Parallel launcher (1 job, 1 node, N tasks)

## Sample “Stress Me!” parameter exploration

Your Turn!

### Hands-on Pre-requisites

► url ◀ | github | src

- Access to ULHPC facility ssh
- (eventually) Open Tmux session (Aion/Iris) `~/.tmux.conf, tmux new -s <name>`  
    ↪ OR (eventually) Open GNU Screen (Iris) `~/.screenrc, screen -S <name>`
- Clone/Pull [ULHPC/tutorials](#) repository `~/git/github.com/ULHPC/tutorials`
- Prepare dedicated directory `~/tutorials/sequential` for this session

```
(access)$> mkdir -p ~/tutorials/sequential
(access)$> cd ~/tutorials/sequential
# create a symbolic link to the reference material
(access)$> ln -s ~/git/github.com/ULHPC/tutorials/sequential/basics ref.d
```

## Sample “Stress Me!” parameter exploration

- **Objective:** mimic (integer) parameter exploration in  $[1 \dots N]$  for serial (1-core) tasks
  - ↳ each task last different amount time (here ranging ranging from 1 to n seconds)
  - ↳ see [scripts/run\\_stressme](#)
    - ✓ stress-based workload generator that imposes a configurable amount of stress on the system

```
./scripts/run_stressme -h
```

### NAME

run\_stressme: A sample workload generator that imposes a stress on the computing node using the 'stress' command.

### USAGE

./scripts/run\_stressme [-n] [N]: run serial stress during N second (Default: 20)

### OPTIONS

-n --noop --dry-run: Dry run mode

## Sample “Stress Me!” parameter exploration

N=#tasks	Expected Seq. time ( $T_1$ )	Optimal time ( $T_\infty$ )
1	1s	1s
10	55s (~1 min)	10s
<b>30 (default)</b>	<b>465s (7 min 45s)</b>	<b>30s</b>
100	5050s (1h 24 min 10s)	100s

### Hands-on: Single task run (interactive)

► [url](#) ◀ | [github](#) | [src](#)

- Get an interactive job
- Open another terminal (or another screen tab/windows)
  - ↪ connect to your running job
  - ↪ monitor it's load with `htop`
- Execute `run_stressme` tasks for various parameters
  - ↪ monitor its effect on second window in `htop`

`si`

`sjoin <JOBID>`  
`htop`



## Sample “Stress Me!” parameter exploration

### Hands-on: A First launcher (1 job, 1 task on one core)

► [url](#) ◀ | [github](#) | [src](#)

- Copy and adapt **ULHPC** Generic serial launcher `launcher.stressme-serial.sh`
- Get an **interactive** job `si`
  - ↪ dry-run vs real execution
- Test **passive** job execution `sbatch`

## Sample “Stress Me!” parameter exploration

- Previous approach works but will **induces an non-necessary overload of the scheduler** when scaling the number of tasks (Ex: job campaign to 1000 or 10000 test cases)
  - ↪ **To complete  $N$  (serial) tasks, you need to submit  $N$  jobs**
    - ✓ **DO NOT make a for loop!!!**: better efficient alternatives **MUST** be used
    - ✓ You will likely wait (sometimes very long time) to have all your jobs completed
  - ↪ Node coverage is sub-optimal
    - ✓ **your serial jobs can be spread on up to  $N$  different nodes**
  - ↪ Very often (90% of the cases), your serial tasks are short-lived (couple of seconds)
    - ✓ You loose more time to schedule jobs than actually running them

## Sample “Stress Me!” parameter exploration

- Previous approach works but will **induces an non-necessary overload of the scheduler** when scaling the number of tasks (Ex: job campaign to 1000 or 10000 test cases)
  - ↪ **To complete N (serial) tasks, you need to submit N jobs**
    - ✓ **DO NOT make a for loop!!!**: better efficient alternatives **MUST** be used
    - ✓ You will likely wait (sometimes very long time) to have all your jobs completed
  - ↪ Node coverage is sub-optimal
    - ✓ **your serial jobs can be spread on up to N different nodes**
  - ↪ Very often (90% of the cases), your serial tasks are short-lived (couple of seconds)
    - ✓ You loose more time to schedule jobs than actually running them
- You **SHOULD** restrict to 1 node **YET** exploit at best the many cores available
  - ↪ Recall: 128 cores per Aion nodes, 28 cores per Iris nodes
  - ↪ Make 1 jobs (for a full nodes) and use it to schedule **inside** your tasks
    - ✓ From scheduler point of view: **only 1 job to schedule** (faster)
    - ✓ **Efficient**: exploit at best the available cores

## Embarrassingly Parallel Jobs

(1 node, 1 job, #cores tasks per node)

- **VERY BAD / To avoid:** Fork [bomb] for loop of sbatch/srun
  - ↳ (dangerous) ampersand & to spawns the command in the background
  - ↳ optimize wait synchronization barrier YET prefer GNU Parallel

```
TASK=run_stressme
ncores=${SLURM_NTASKS_PER_NODE:-$(nproc --all)} # Number of cores per node
for i in {1..1000}; do # DON'T DO THAT - just to illustrate
    srun -nl --exclusive -c 1 --cpu-bind=cores ${TASK} $i & # NO!! Fork 1000 sub-processes
done
wait
```

## Embarrassingly Parallel Jobs

(1 node, 1 job, #cores tasks per node)

- **VERY BAD / To avoid:** Fork [bomb] for loop of sbatch/srun
  - ↳ (dangerous) ampersand & to spawns the command in the background
  - ↳ optimize wait synchronization barrier YET prefer GNU Parallel

```
TASK=run_stressme
ncores=${SLURM_NTASKS_PER_NODE:-$(nproc --all)} # Number of cores per node
for i in {1..1000}; do # DON'T DO THAT - just to illustrate
    srun -nl --exclusive -c 1 --cpu-bind=cores ${TASK} $i &
    [[ $((i%ncores)) -eq 0 ]] && wait # Better but suboptimal
done
wait
```

## Embarrassingly Parallel Jobs

(1 node, 1 job, #cores tasks per node)

- **VERY BAD / To avoid:** Fork [bomb] for loop of sbatch/srun
  - ↳ (dangerous) ampersand & to spawns the command in the background
  - ↳ optimize wait synchronization barrier YET prefer GNU Parallel
- **BAD / To avoid:** Job arrays `sbatch --array 1-1000%128 [...]`
  - ↳ in all cases, job array capabilities are drastically reduced (on purpose)
    - ✓ `scontrol show config | grep -i maxarray`

Job Array Variable	Description
<code>\$SLURM_ARRAY_JOB_ID</code>	first job ID of the array
<code>\$SLURM_ARRAY_TASK_ID</code>	current job array step index value.
<code>\$SLURM_ARRAY_TASK_COUNT</code>	number of tasks in the job array.

## Embarrassingly Parallel Jobs

(1 node, 1 job, #cores tasks per node)

- **VERY BAD / To avoid:** Fork [bomb] for loop of sbatch/srun
  - ↳ (dangerous) ampersand & to spawns the command in the background
  - ↳ optimize wait synchronization barrier YET prefer GNU Parallel
- **BAD / To avoid:** Job arrays sbatch --array 1-1000%128 [...]
  - ↳ in all cases, job array capabilities are drastically reduced (on purpose)
    - ✓ `scontrol show config | grep -i maxarray`

Job Array Variable	Description
<code>\$SLURM_ARRAY_JOB_ID</code>	first job ID of the array
<code>\$SLURM_ARRAY_TASK_ID</code>	current job array step index value.
<code>\$SLURM_ARRAY_TASK_COUNT</code>	number of tasks in the job array.

`${TASK} ${SLURM_ARRAY_TASK_ID} # Simpler syntax YET induce same (worse) load on scheduler`

## Embarrassingly Parallel Jobs

(1 node, 1 job, #cores tasks per node)

- **VERY BAD / To avoid:** Fork [bomb] for loop of sbatch/srun
  - ↳ (dangerous) ampersand & to spawns the command in the background
  - ↳ optimize wait synchronization barrier YET prefer GNU Parallel
- **BAD / To avoid:** Job arrays sbatch --array 1-1000%128 [...]
  - ↳ in all cases, job array capabilities are drastically reduced (on purpose)
    - ✓ `scontrol show config | grep -i maxarray`

Job Array Variable	Description
\$SLURM_ARRAY_JOB_ID	first job ID of the array
\$SLURM_ARRAY_TASK_ID	current job array step index value.
\$SLURM_ARRAY_TASK_COUNT	number of tasks in the job array.

- **Always prefer GNU Parallel for Embarrassingly Parallel Serial Tasks!**
  - ↳ **Good, efficient and Optimized**
  - ↳ Transparent through the provided ULHPC launcher





# Summary

- 1 Introduction  
[Serial] Task, Job and scheduling considerations  
Slurm Launchers Templates for Sequential Applications
- 2 A First launcher (1 job, 1 task on one core)
- 3 Embarrassingly parallel jobs with GNU Parallel launcher (1 job, 1 node, N tasks)**

## GNU Parallel



- A shell tool to execute **independent** tasks in parallel using one or more nodes
  - ↪ A task can be a single command or script
- Typical input could be a list of files, a list of parameters
  - ↪ Recall that in bash, "{1..10}" is expanded to "1 2 3 4 5 6 7 8 9 10"

### Hands-on: GNU Parallel Build and Installation

► url ◀ | github | src

- Setup your HOME to support GNU Stow installation
  - ↪ GNU Stow manual - tutorial
- Download, build and install parallel
- Use stow to enable/disable the installed version

stow [-D]

On stow -D, you can ignore the 'BUG in find\_stowed\_path?' message due to recent [Issues #65](#)

## GNU Parallel Syntax

GNU Parallel syntax can be a little distributing, but basically it supports two modes:

- 1 Reading command arguments **on the command line**:

```
parallel [-j N] [OPTIONS] COMMAND {} ::: TASKLIST
```

- 2 Reading command arguments **from an input file**:

```
parallel -a TASKLIST.LST [-j N] [OPTIONS] COMMAND {}  
parallel [-j N] [OPTIONS] COMMAND {} :::: TASKLIST.LST
```

- **Tips & Tricks:**

- If your COMMAND embed a pipe stage, you have to escape the pipe symbol as follows \
- **ALWAYS** echo your commands until you're satisfied



Embarrassingly parallel jobs with GNU Parallel launcher (1 job, 1 node, N tasks)

## GNU Parallel Syntax

### Hands-on: Discovering the parallel command

► [url](#) ◀ | [github](#) | [src](#)

- play with a TASKLIST from the command line
- play now with a TASKLIST from an input file.

## GNU Parallel Syntax

```
# Default string {} --> path/to/filename.ext
parallel echo {} ::: path/to/filename.ext
# Remove extension {.} --> path/to/filename
parallel echo {.} ::: path/to/filename.ext
# Remove path {/} --> filename.ext
parallel echo {/} ::: path/to/filename.ext
# Remove path and extension {/.} --> filename
parallel echo {/.} ::: path/to/filename.ext
# Change extension and path --> output/filename.out
parallel echo output/{/.}.out ::: path/to/filename.ext
```

## Back to “Stress Me!” parameter exploration

### Hands-on: Best launcher based on GNU Parallel

► url ◄ | github | src

- Copy and adapt **ULHPC Generic GNU Parallel launcher**
  - `launcher.stressme.${ULHPC_CLUSTER}.sh`
  - only difference optimized default `--ntasks-per-node` for aion/iris
  - convenient symlink `launcher.stressme.sh`
- Get an **exclusive interactive** job `si --exclusive`
  - dry-run vs test vs real execution
- Test **passive** job execution `sbatch`
- Embarrassingly [GNU] parallel tasks across multiples nodes
  - Concepts and Safeguards
  - Slurm dependency mechanism `--dependency=singleton`
  - sample submission script `scripts/submit_stressme_multinode`

## Back to “Stress Me!” parameter exploration

```
(access)$> ./scripts/submit_stressme_multinode -h
Usage: submit_stressme_multinode [-x] [-N MAXNODES]
    Sample submission script across multiple nodes
    Execution won't spread on more than 4 nodes (singleton dependency)
    -x --execute          really submit the jobs with sbatch
    -N --nodes MAXNODES  set max. nodes
# Target restriction to 2 running nodes max
(access)$> ./scripts/submit_stressme_multinode -N 2
(access)$> ./scripts/submit_stressme_multinode -N 2 -x
# queue -u $(whoami)
  JOBID PARTIT  QOS      NAME  NODE  CPUS ST  TIME  TIME_LEFT NODELIST(REASON)
  2175780 batch normal StressMe_0    1   28 PD  0:00   1:00:00 (Dependency)
  2175781 batch normal StressMe_1    1   28 PD  0:00   1:00:00 (Dependency)
  2175782 batch normal StressMe_0    1   28 PD  0:00   1:00:00 (Dependency)
  2175779 batch normal StressMe_1    1   28  R  0:02   59:58 iris-064
  2175778 batch normal StressMe_0    1   28  R  0:05   59:55 iris-047
```



Thank you for your attention...

## Questions?

[ulhpc-tutorials.rtf.d.io/en/latest/sequential/basics](http://ulhpc-tutorials.rtf.d.io/en/latest/sequential/basics)



### High Performance Computing @ Uni.lu

University of Luxembourg, Belval Campus  
Maison du Nombre, 4th floor  
2, avenue de l'Université  
L-4365 Esch-sur-Alzette  
mail: [hpc@uni.lu](mailto:hpc@uni.lu)

1

#### Introduction

[Serial] Task, Job and scheduling considerations  
Slurm Launchers Templates for Sequential Applications

2

#### A First launcher (1 job, 1 task on one core)

3

#### Embarrassingly parallel jobs with GNU Parallel launcher (1 job, 1 node, N tasks)

## Uni.lu HPC School 2021 Contributors

	<b>Dr. Xavier Bessoron</b> Research Scientist		<b>Abatcha Ollah</b> Infra. & HPC Arch. Engineer
	<b>Hyacinthe Cartiaux</b> Infra. & HPC Arch. Engineer		<b>Dr. Tiago C. Pessoa</b> Postdoctoral Researcher
	<b>Dr. Aurelien Ginohac</b> Research Scientist		<b>Sarah Peter</b> Infra. & Arch. Engineer
	<b>Dr. Emmanuel Kieffer</b> Research Scientist		<b>Teddy Valette</b> Infra. & HPC Arch. Engineer
	<b>Dr. Loizos Koutsantonis</b> Postdoctoral Researcher		<b>Dr. Sebastien Varrette</b> Research Scientist
	<b>Dr. Ezhilmathi Krishnasamy</b> Postdoctoral Researcher	... and additional help (Survey, session tests)	
			<b>Arlyne Vandeventer</b> Project Manager

[hpc.uni.lu](http://hpc.uni.lu)